

Versioning and Experiment Tracking



Why do we need experiment tracking tools?

At first glance, building and deploying machine learning models looks a lot like writing code.

But there are some key differences that make machine learning harder:

- Have far more branching and experimentation than a typical software project.
- Rather than throwing errors, it just underperforms, making debugging extra difficult and time consuming.
- A single small change in training data, training code or hyperparameters can wildly change a model's performance, so reproducing earlier work often requires exactly matching the prior setup.
- Running machine learning experiments can be time consuming and compute costs may be expensive.

















PROBLEM

Massive developer tools gap for ML

SOFTWARE 1.0 - WRITE CODE

Design	Version Code	Code & Collaborate	Deploy to Infra	CI/CD	Production Monitoring
 Figma  Sketch	 GitHub  GitLab	 Jira  VS Code	 HashiCorp  puppet	 circleci  Jenkins	 PagerDuty  DATADOG

SOFTWARE 2.0 - TRAIN MODELS

Prep & Visualize Data	Version Data & Models	Experiment Tracking	Manage Model Pipeline	Model CI/CD	Production monitoring
Custom apps Notebooks	Files in S3	Text files Screenshots	Text files	Custom scripts	Nothing

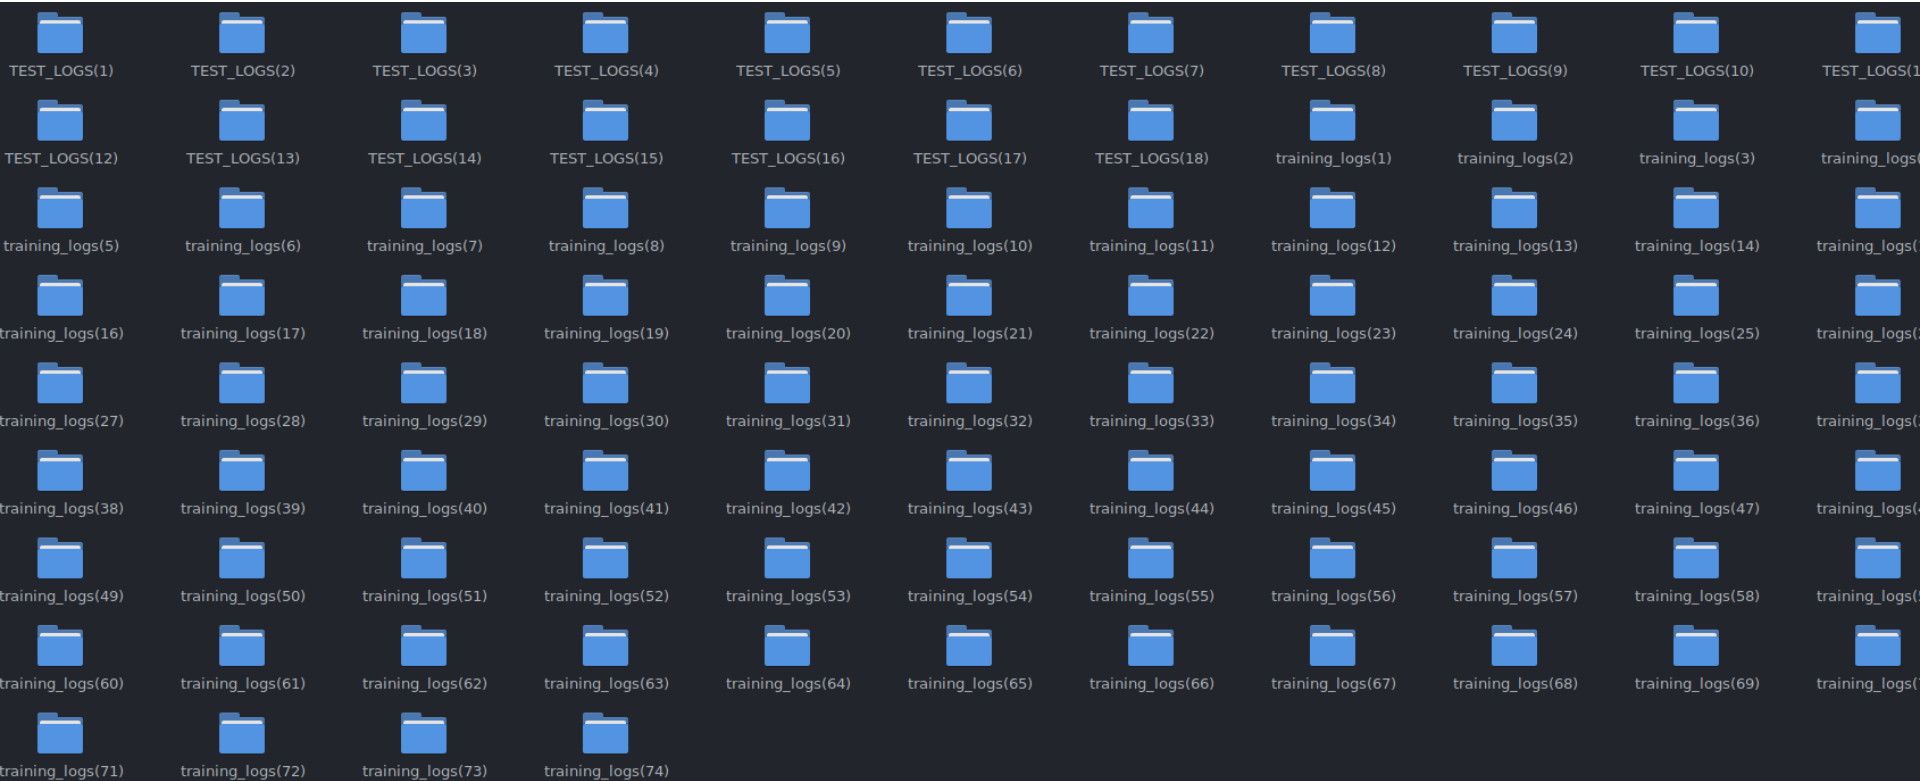


Why do we need experiment tracking tools?

Suppose you are training a model with different options on;

- Training-validation-test split
- Device (RTX2080Ti, TPU v2 pod, etc.)
- Hyperparameters
- Training procedure (eg. FPN for object detection)
- ...





Why do we need experiment tracking tools?

These tools take away the burden of;

- Manually managing experiment runs
- Finding logs of specific hyper-parameters or configurations.
- Managing version of dataset (optional)
- Managing orchestration (optional)



	Neptune	Weights & Biases	Comet	Sacred + Omniboard	MLflow	TensorBoard	Guild AI	Polyaxon	TRAINS	Valohai	Pachyderm	Kubeflow	Verta.ai	SageMaker Studio	DVC
Overview															
Focus	Experiment Management	Experiment Management	Experiment Management	Experiment Management	Entire Lifecycle	Experiment Management	Experiment Management	Experiment Management	Experiment Management	Entire Lifecycle	Entire Lifecycle	Run orchestration	Entire Lifecycle	Entire Lifecycle	Data Versioning
Price	<ul style="list-style-type: none"> Free Team Trial: 1 project Team Research: \$0 Team Startup: \$39 per user Team Pro: \$79 per user Enterprise: starts at \$1799 	<ul style="list-style-type: none"> Personal Startup: \$35 per user Team: \$100 per user Enterprise: NA 	<ul style="list-style-type: none"> Free Startup: \$39 per user Teams: \$179 per user Enterprise: NA 	Free	Free	Free	Free	Enterprise: NA	Free	<ul style="list-style-type: none"> Free Pro: \$649 Enterprise: NA 	<ul style="list-style-type: none"> Community Edition Enterprise Edition: NA 	Free	Enterprise: NA	You pay extra on top of compute	Free
Free plan limitation	<ul style="list-style-type: none"> Free: 1 user Team Free: 1 project Team Research: no limits 	<ul style="list-style-type: none"> 1 user limited support 	<ul style="list-style-type: none"> 1 user no private links (sharing) 							only 3 days of data available					
Open – source				✓	✓	✓	✓	limited	✓		limited	✓			✓
Easy integration	✓	✓	✓	<ul style="list-style-type: none"> Somewhat difficult 	✓	✓	✓								✓
				<ul style="list-style-type: none"> Depends on the on backend 											
Lightweight	✓	✓	✓	Somewhat	✓	✓	✓		✓						✓

source: <https://neptune.ai/wp-content/uploads/best-machine-learning-experiment-tracking-apps.pdf>



Benefits over Tensorboard

Cloud storage of results

- Remote tracking of experiments.

Hyperparameter Tuning

- Self contained or integration with other tools.

Easier tracking

- Better UI
- Comprehensible statistics over hardware usage and parameter distributions

Framework Agnostic

- Can work with any framework (Even works without any framework)



Weights & Biases

Any Framework

TensorFlow

PyTorch

Keras

Scikit

Hugging Face

XGBoost

```
# Flexible integration for any Python script
import wandb

# 1. Start a W&B run
wandb.init(project='gpt3')

# 2. Save model inputs and hyperparameters
config = wandb.config
config.learning_rate = 0.01

# Model training here

# 3. Log metrics over time to visualize performance
wandb.log({"loss": loss})
```



Weights & Biases


Any Framework	import wandb
TensorFlow	
PyTorch	# 1. Start a new run wandb.init(project="gpt-3")
Keras	
Scikit	# 2. Save model inputs and hyperparameters config = wandb.config config.learning_rate = 0.01
Hugging Face	
XGBoost	# 3. Log gradients and model parameters wandb.watch(model) for batch_idx, (data, target) in enumerate(train_loader): ... if batch_idx % args.log_interval == 0: # 4. Log metrics to visualize performance wandb.log({"loss": loss})




ML Developer Tools for the entire ML workflow



W&B PLATFORM



FRAMEWORK AGNOSTIC



ENVIRONMENT AGNOSTIC



Weights & Biases

1. Install the lightweight library
2. Track metrics and artifacts
3. Compare experiments fast
4. Optimize hyperparameters easily
5. Collaborate on reproducible models



Light integration

Try W&B in 5 minutes:

- Intro notebook
bit.ly/intro-wb
- Quickstart docs
docs.wandb.com

PYTORCH

```
import torch  
  
model.train()
```

```
import torch  
import wandb  
  
wandb.watch(model)  
model.train()
```

TENSORFLOW

```
import tensorflow as tf  
  
classifier.train()
```

```
import tensorflow as tf  
import wandb  
  
wandb.init(sync_tensorboard=True)  
  
classifier.train()
```

KERAS

```
from tensorflow import keras  
  
model.fit(X, y)
```

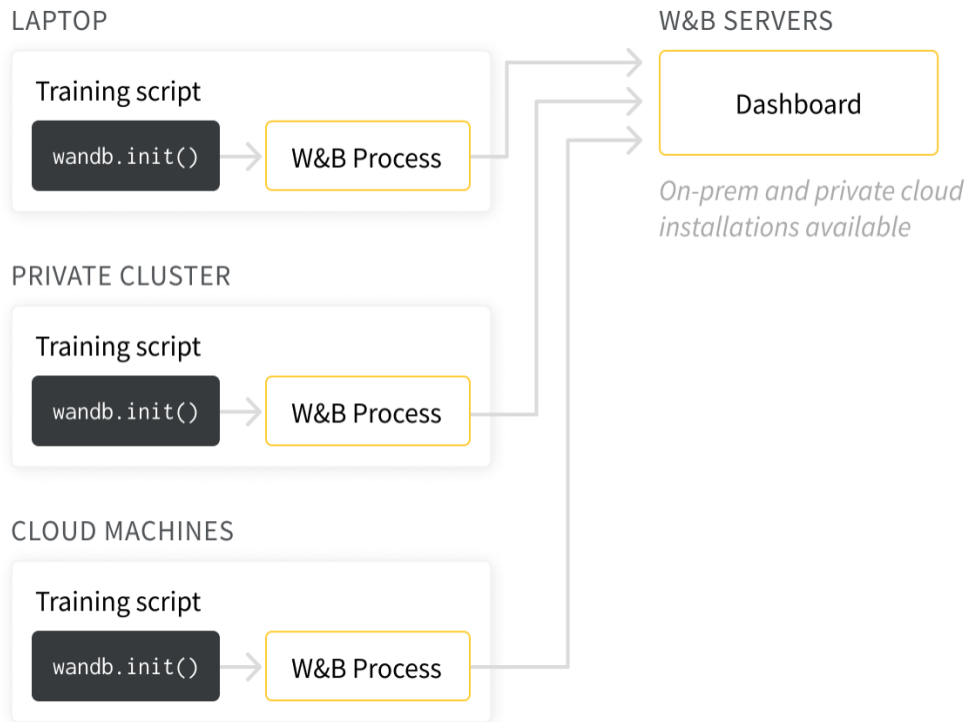
```
from tensorflow import keras  
import wandb  
  
model.fit(X, y,  
          callbacks=[WandbCallback()])
```



Infra agnostic

W&B can be used


- In their hosted cloud
- In a private cloud
- On-premise-on your own machines



Track experiments

See live updates on model performance, check for overfitting, and visualize how a model performs on different classes.

bit.ly/demo-run



Track experiments – Before W&B

Layer (type)	Output Shape	Param #	Connected to
conv2d_1_input (InputLayer)	(None, 299, 299, 3)	0	
lambda_1 (Lambda)	(None, 299, 299, 3)	0	conv2d_1_input[0][0]
lambda_2 (Lambda)	(None, 299, 299, 3)	0	conv2d_1_input[0][0]
sequential_1 (Sequential)	(None, 10)	910922	lambda_1[0][0] lambda_2[0][0]
activation_7 (Concatenate)	(None, 10)	0	sequential_1[1][0] sequential_1[2][0]
Total params: 910,922			
Trainable params: 910,922			
Non-trainable params: 0			
None			
Found 49999 images belonging to 10 classes.			
Found 8000 images belonging to 10 classes.			
Epoch 1/50			
39/39 [=====] - 270s 7s/step - loss: 2.3153 - acc: 0.1178 - val_loss: 2.2428 - val_acc: 0.1589			
Epoch 2/50			
39/39 [=====] - 263s 7s/step - loss: 2.2588 - acc: 0.1538 - val_loss: 2.1890 - val_acc: 0.2174			
Epoch 3/50			
39/39 [=====] - 262s 7s/step - loss: 2.2060 - acc: 0.1827 - val_loss: 2.1767 - val_acc: 0.2096			
Epoch 4/50			
39/39 [=====] - 263s 7s/step - loss: 2.1640 - acc: 0.2107 - val_loss: 2.1133 - val_acc: 0.2357			
Epoch 5/50			
39/39 [=====] - 261s 7s/step - loss: 2.0827 - acc: 0.2432 - val_loss: 2.1499 - val_acc: 0.2344			
Epoch 6/50			
39/39 [=====] - 262s 7s/step - loss: 2.0810 - acc: 0.2508 - val_loss: 2.0289 - val_acc: 0.2604			
Epoch 7/50			
39/39 [=====] - 262s 7s/step - loss: 2.0575 - acc: 0.2602 - val_loss: 1.9912 - val_acc: 0.2865			
Epoch 8/50			
39/39 [=====] - 261s 7s/step - loss: 2.0355 - acc: 0.2734 - val_loss: 2.0319 - val_acc: 0.2409			
Epoch 9/50			
39/39 [=====] - 263s 7s/step - loss: 2.0254 - acc: 0.2829 - val_loss: 1.9364 - val_acc: 0.3307			
Epoch 10/50			
39/39 [=====] - 262s 7s/step - loss: 2.0056 - acc: 0.2804 - val_loss: 1.9934 - val_acc: 0.2773			
Epoch 11/50			
39/39 [=====] - 256s 7s/step - loss: 1.9678 - acc: 0.3048 - val_loss: 1.9048 - val_acc: 0.3352			
Epoch 12/50			
39/39 [=====] - 260s 7s/step - loss: 1.9634 - acc: 0.3109 - val_loss: 1.8758 - val_acc: 0.3568			

Table 3: Detection results on **PASCAL VOC 2007 test set**. The detector is Fast R-CNN and VGG-16. Training data: “07”: VOC 2007 trainval, “07+12”: union set of VOC 2007 trainval and VOC 2012 trainval. For RPN, the train-time proposals for Fast R-CNN are 2000. [†]: this number was reported in [2]; using the repository provided by this paper, this result is higher (68.1).

method	# proposals	data	mAP (%)
SS	2000	07	66.9 [†]
SS	2000	07+12	70.0
RPN+VGG, unshared	300	07	68.5
RPN+VGG, shared	300	07	69.9
RPN+VGG, shared	300	07+12	73.2
RPN+VGG, shared	300	COCO+07+12	78.8

Table 4: Detection results on **PASCAL VOC 2012 test set**. The detector is Fast R-CNN and VGG-16. Training data: “07”: VOC 2007 trainval, “07+12”: union set of VOC 2007 trainval+test and VOC 2012 trainval. For RPN, the train-time proposals for Fast R-CNN are 2000. [†]: <http://host.robots.ox.ac.uk:8080/anonymous/HZJTQA.html>. [‡]: <http://host.robots.ox.ac.uk:8080/anonymous/YNPLXB.html>. [§]: <http://host.robots.ox.ac.uk:8080/anonymous/XEDH10.html>.

method	# proposals	data	mAP (%)
SS	2000	12	65.7
SS	2000	07+12	68.4
RPN+VGG, shared [†]	300	12	67.0
RPN+VGG, shared [†]	300	07+12	70.4
RPN+VGG, shared [†]	300	COCO+07+12	75.9

Table 5: **Timing** (ms) on a K40 GPU, except SS proposal is evaluated in a CPU. “Region-wise” includes NMS, pooling, fully-connected, and softmax layers. See our released code for the profiling of running time.

model	system	conv	proposal	region-wise	total	rate
VGG	SS + Fast R-CNN	146	1510	174	1830	0.5 fps
VGG	RPN + Fast R-CNN	141	10	47	198	5 fps
ZF	RPN + Fast R-CNN	31	3	25	59	17 fps

	backbone	AP	AP ₅₀	AP ₇₅	AP _S	AP _M	AP _L
<i>Two-stage methods</i>							
Faster R-CNN+++ [5]	ResNet-101-C4	34.9	55.7	37.4	15.6	38.7	50.9
Faster R-CNN w FPN [8]	ResNet-101-FPN	36.2	59.1	39.0	18.2	39.0	48.2
Faster R-CNN by G-RMI [6]	Inception-ResNet-v2 [21]	34.7	55.5	36.7	13.5	38.1	52.0
Faster R-CNN w TDM [20]	Inception-ResNet-v2-TDM	36.8	57.7	39.2	16.2	39.8	52.1
<i>One-stage methods</i>							
YOLOv2 [15]	DarkNet-19 [15]	21.6	44.0	19.2	5.0	22.4	35.5
SSD513 [11, 3]	ResNet-101-SSD	31.2	50.4	33.3	10.2	34.5	49.8
DSSD513 [3]	ResNet-101-DSSD	33.2	53.3	35.2	13.0	35.4	51.1
RetinaNet [9]	ResNet-101-FPN	39.1	59.1	42.3	21.8	42.7	59.9
RetinaNet [9]	ResNeXt-101-FPN	40.8	61.1	44.1	24.1	44.2	51.2
YOLOv3 608 × 608	Darknet-53	33.0	57.9	34.4	18.3	35.4	41.9

Track experiments – After W&B



batch 64 local

What makes this run special?

Privacy: PUBLIC

Tags: 2GPU, keras, working

Author: stacey

State: finished

Start time: March 22nd, 2019 at 7:28:49 am

Duration: 2h 47m 8s

Run path: stacey/estuary/kg/zuc6y

Hostname: wbeast

OS: Linux-4.15.0-46-generic-x86_64-with-debian-stretch-sid

Python version: 2.7.13

Git repository: git clone git+ssh://git@github.com:wandb/explored1.git

Git state: git checkout -b "batch-64-local" 9b4586dfb3a265de98ad6c85ec8be778567289

Command: experiment_staging.py -m distrib_keras_util_multi_gpu --distrib -b 64 -nt 6400 -nv 1280

System Hardware: CPU count 12, GPU count 2, GPU type GeForce GTX 1080 Ti

Config

Config parameters describe your model's inputs. [Learn more](#)

Name	Value
batch_size	64
dropout	0.3
epochs	50
img_dim	299
n_conv_layers	1
n_fc_layers	0
n_train	6400

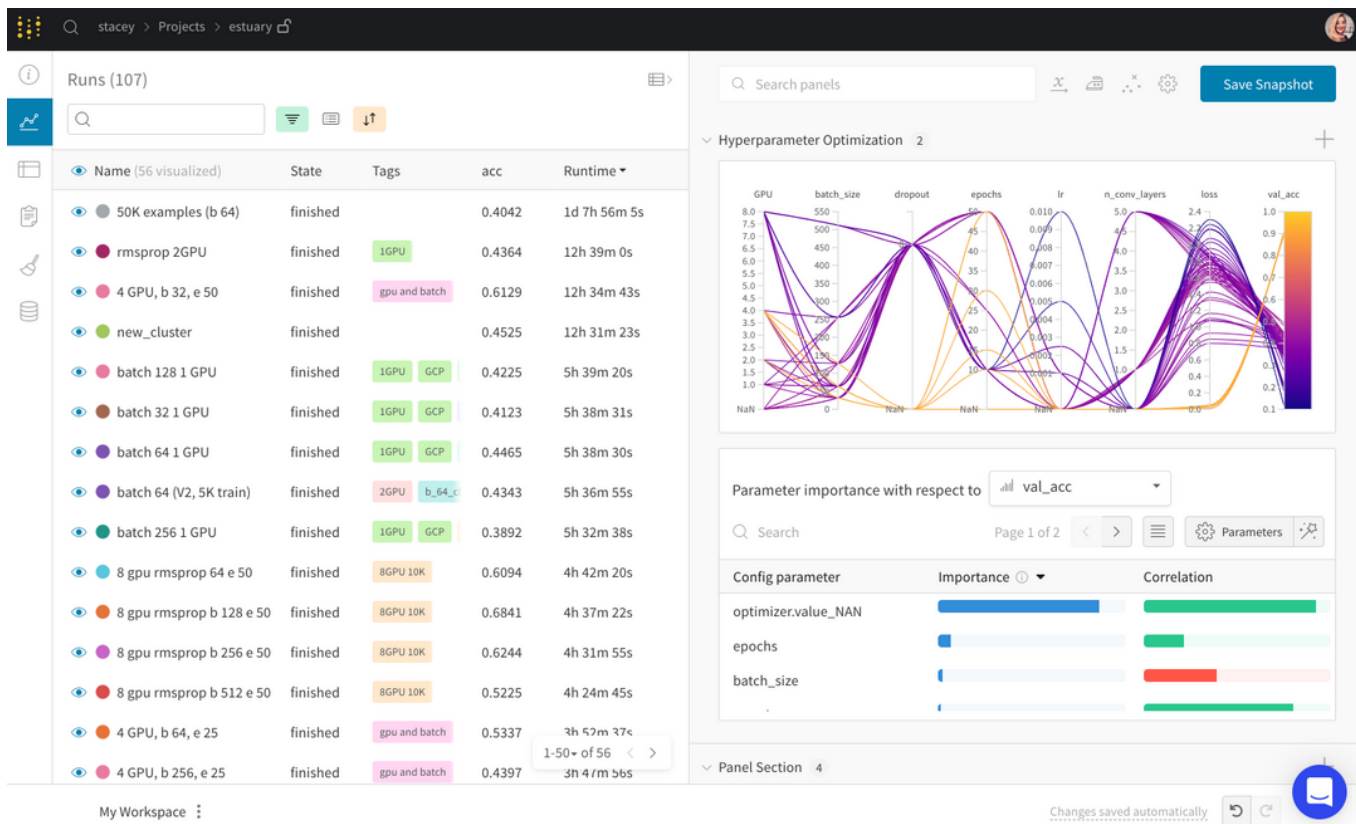
Summary

Summary metrics describe your results.

Name	Value
acc	
epoch	
graph	
loss	
val_acc	
val_loss	



Track experiments – After W&B

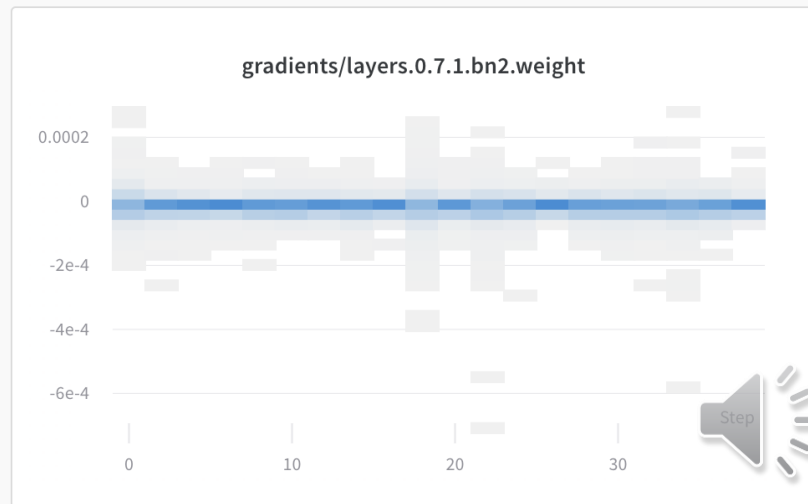



[Live Dashboard](#)

bit.ly/demo-run

Visualize gradients



See the inner workings of your model, look for convergence during training, and check for exploding gradients.



System metrics

Get the most out of your GPUs
and identify opportunities for
optimizing hardware utilization.

bit.ly/demo-run



Capture the code

Save the most recent git commit, the command and args, and system hardware setup.

W&B also saves a patch file with uncommitted changes so you can reproduce the exact code that trained the model.

bit.ly/demo-run

dutiful-dragon-174 

What makes this run special? 

Privacy

 PUBLIC

Tags



Author

stacey

State

finished

Start time

January 24th, 2020 at 1:30:16 pm

Duration

38m 27s

Run path

stacey/deep-drive/6zsn8ltb

Hostname

wbrave

OS

Linux-5.0.0-37-generic-x86_64-with-debian-buster-sid

Python version

3.7.2

Python executable `/home/stacey/.pyenv/versions/sd/bin/python`

Git repository

`git clone https://github.com/borisdayma/semantic-segmentation.git`

Git state

`git checkout -b "dutiful-dragon-174" f0a9494a5d152663d226e28e279c65`

Command

`train.py`

CPU count 20

System Hardware

GPU count 2

GPU type GeForce RTX 2080 Ti

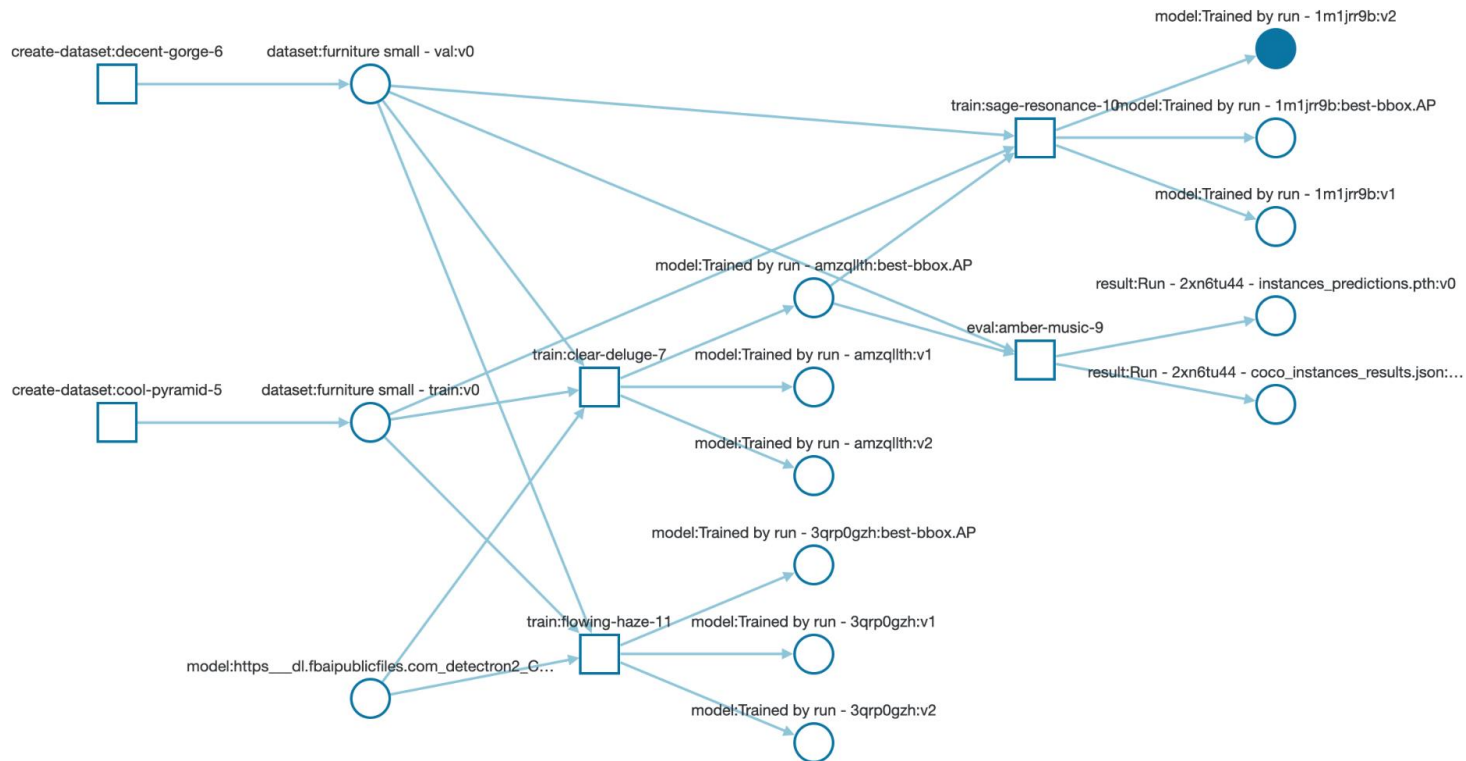
W&B CLI Version 0.8.21



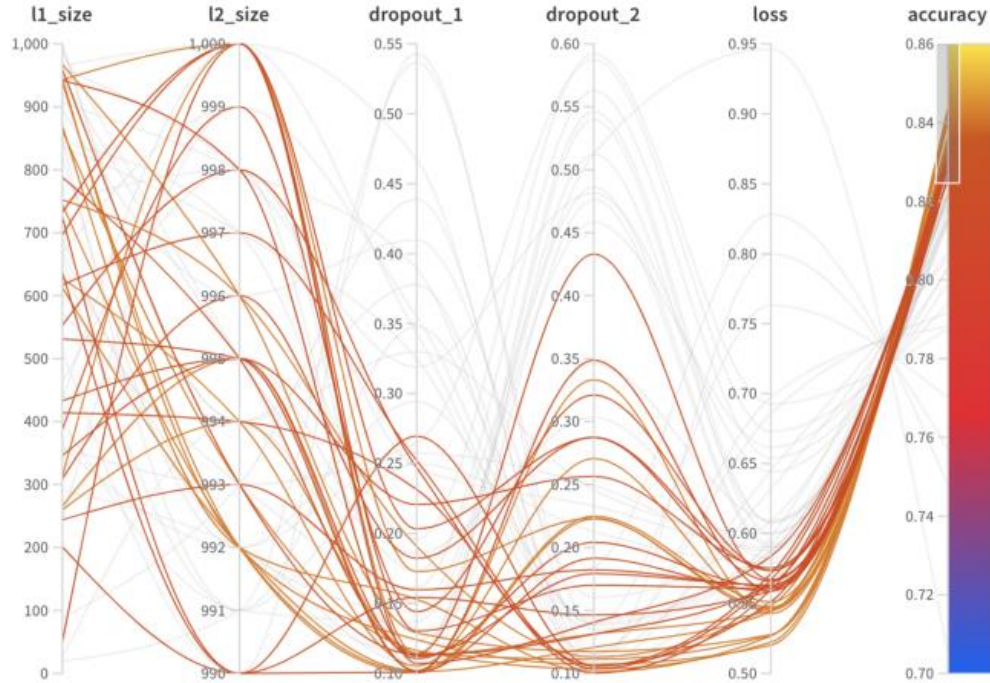
Track artifacts



Track artifacts



Compare runs

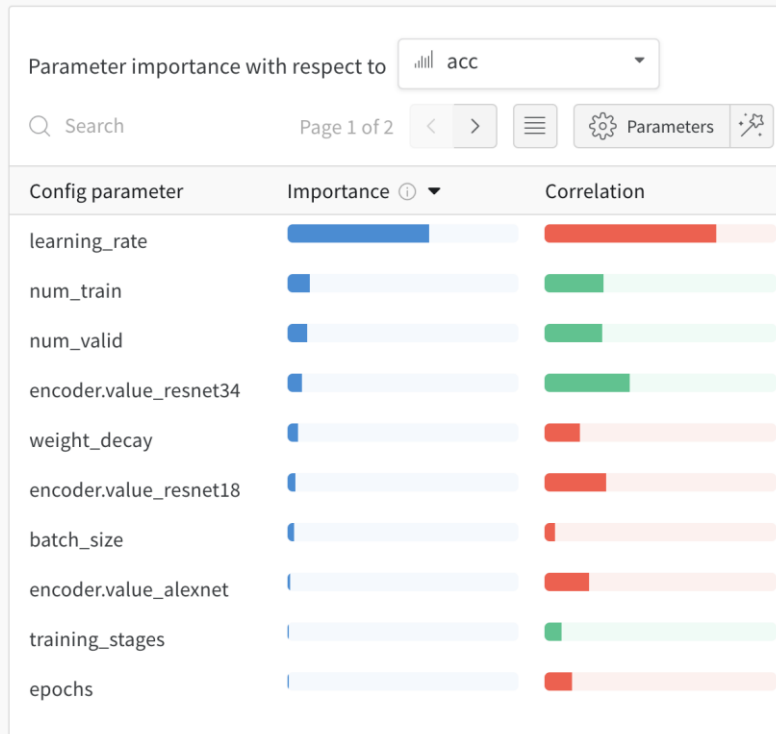


Compare runs

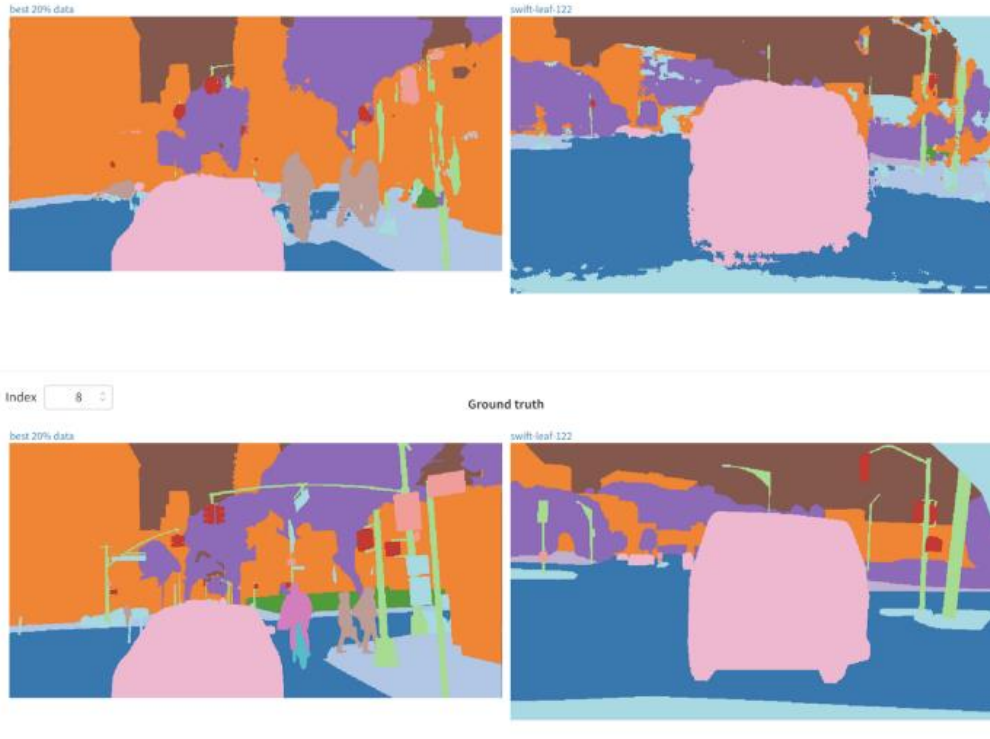
Visualize the relationships between hyperparameters and model metrics.

Explore the space of possible models quickly, without getting bogged down setting up manual visualizations.

bit.ly/deep-drive



Viewing Specific Examples



Persistent system of record

Query and filter across runs to keep projects organized.

Runs (395)



Search 3 Filters Group Sort Tag Move Columns

<input type="checkbox"/> Name (228 visualized)	Tags	Runtime	batch_size	encoder	learning_rate	num_train	num_valid	weight_decay	iou	train_loss	valid_loss	acc	traffic_acc	road
best car acc (50% data)	seg_masks	47m 52s	6	resnet34	0.001311	3524	492	0.08173	0.7997	0.5375	0.4427	0.8823	0.8664	0.93
best traffic acc (50% data)	seg_masks	46m 42s	8	resnet18	0.001	3523	492	0.097	0.8073	0.4919	0.4203	0.888	0.8718	0.94
best human iou (50% data)	seg_masks	31m 34s	7	alexnet	0.0009084	1405	190	0.097	0.716	0.6222	0.6259	0.8334	0.8042	0.9
best overall IOU (20% data)	seg_masks	20m 23s	7	resnet34	0.001367	1376	205	0.06731	0.7948	0.5096	0.4659	0.8726	0.8593	0.93
vibrant-cherry-426		6m 12s	8	resnet34	0.001	347	42	0.097	0.752	0.7114	0.6055	0.8474	0.8282	0.95
worldly-totem-422		12m 54s	8	resnet34	0.001	682	97	0.097	0.7523	0.5774	0.5836	0.8566	0.8349	0.91
jumping-voice-421		11m 59s	8	resnet34	0.001	725	92	0.097	0.7449	0.5633	0.5334	0.8504	0.8296	0.91
logical-energy-420	test_only	2m 14s	8	resnet34	0.001	66	10	0.097	0.4297	1.459	1.221	0.626	0.5958	0.76



Hyperparameter Sweeping

- Hyperparameter sweep is **a way of finding the best hyperparameters for a model, given a set of data.**
- This is done by automatically searching through combinations of hyperparameter values (e.g. learning rate, batch size, number of hidden layers, optimizer type) to find the most optimal values.



Hyperparameter Sweeping

1. **Define the sweep:** Create a [YAML file](#) that specifies the parameters to search through, the search strategy and the optimization metric.
2. **Initialize the sweep:** Initialize the sweep and pass in the dictionary of sweep configurations: `"sweep_id = wandb.sweep(sweep_config)"`
3. **Run the sweep agent:** Call `wandb.agent()` and pass the `sweep_id` to run, along with a function that defines your model architecture and trains it: `"wandb.agent(sweep_id, function=train)"`



Hyperparameter Sweeping

method	Description
grid	Grid search iterates over all possible combinations of parameter values.
random	Random search chooses a random set of values on each iteration.
bayes	Our Bayesian hyperparameter search method uses a Gaussian Process to model the relationship between the parameters and the model metric and chooses parameters to optimize the probability of improvement. This strategy requires the <code>metric</code> key to be specified.



Collaborative Reports

Reports let you organize and embed visualizations, describe your findings, share updates with collaborators.

Notes: Add a graph with a quick note to yourself.


Collaboration: Share findings with your colleagues.

Work log: Track what you've tried and plan next steps

bit.ly/deep-drive-report

Inception V3 parallelizes better than Inception-ResNet-V2

Plotting my GPU usage across the two models explains what's happening. The right plot below shows the GPU utilization percentage for each of 8 GPUs in red colors for the Red 2016 Inception-ResNet V2 model and in blue colors for Blue 2015 Inception V3. In the Red 2016 model, GPU 0 (top orange line) does way more work: 6-8X the work of the other 7 GPUs (dropping at the end of the first epoch, at around 46 minutes). In the Blue 2015 model, all 8 GPUs share work much more evenly. Using Inception V3 instead of Inception-ResNet-V2 for this task will let me iterate much faster.



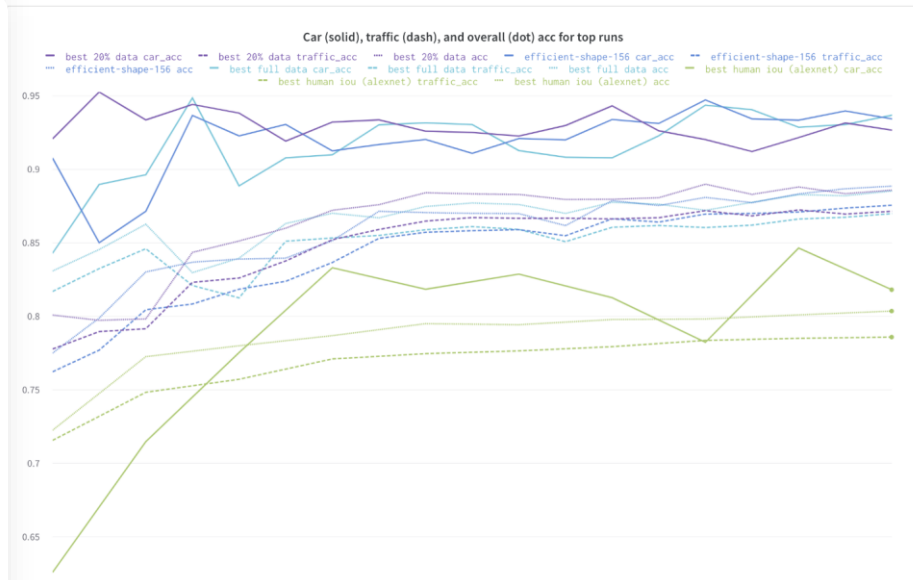
REPORTS

Live dashboards

Customize real-time views of model training and evaluation, and share auto-updating dashboards with your team to keep everyone in the loop.

bit.ly/deep-drive-report

Comparing per-class accuracies



☒ Four top models 4

Name (4 visualized)	acc ▼	car_acc	traffic_acc	train_loss	valid_loss	encoder	learning_ra	num_train	num_val
best 20% data	0.8866	0.931	0.8724	0.4421	0.4284	resnet18	0.001	1421	183
efficient-shape-156	0.8863	0.9329	0.8739	0.4321	0.4394	resnet18	0.001	1409	207
best full data	0.8848	0.9356	0.8697	0.5218	0.4242	resnet18	0.001	700	100
best human iou (alexnet)	0.8036	0.8181	0.7859	0.6373	0.708	alexnet	0.0009084	699	100

1-4 of 4 < >

Experiment Tracking In Practice

OpenAI

- Reports
 - Every experiment the OpenAI team tries starts as a Report
 - “The ability to mix real data from experiments with context and commentary on the results was the primary selling point for us”
 - From Google docs with screenshots to interactive reports.
- Workflow
 - Create a new report with text at the top explaining the context, hypotheses under test, and experimental plan.
 - Launch one or more experiments aimed at testing the hypotheses.
 - Monitor the experiments. If enough data has been collected after the first set of experiments, add a conclusion to the top and share with the team. Otherwise, update the hypotheses and experimental plan and return to step 2.

<https://wandb.ai/openai/published-work/Learning-Dexterity-End-to-End--VmIld7oxMTUvMDQ>

